

Start coding or [generate](#) with AI.

Double-click (or enter) to edit





Task 4: Named Entity Recognition (NER) from News Articles

Description:

- Dataset (Recommended): CoNLL003 (Kaggle)
- Identify named entities (like people, locations, and organizations) from article content
- Use rule-based and model-based NER approaches
- Highlight and categorize extracted entities in the text

✓ BERT-based Named Entity Recognition (NER) model

```
!pip install seqeval==1.2.2
```



Collecting seqeval==1.2.2

Downloading seqeval-1.2.2.tar.gz (43 kB)

43.6/43.6 kB 2.3 MB/s eta 0:00:00

Preparing metadata (setup.py) ... done

Requirement already satisfied: numpy>=1.14.0 in /usr/local/lib/python3.12/dist-packages

Requirement already satisfied: scikit-learn>=0.21.3 in /usr/local/lib/python3.12/dist-packages

```
Requirement already satisfied: scipy>=1.6.0 in /usr/local/lib/python3.12/dist-packages (
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.12/dist-pa
Building wheels for collected packages: sequeval
  Building wheel for sequeval (setup.py) ... done
  Created wheel for sequeval: filename=sequeval-1.2.2-py3-none-any.whl size=16162 sha256=a
  Stored in directory: /root/.cache/pip/wheels/5f/b8/73/0b2c1a76b701a677653dd79ece07cfab
Successfully built sequeval
Installing collected packages: sequeval
Successfully installed sequeval-1.2.2
```

This dataset is divided into train.txt, test.txt and valid.txt The Tokens are labeled under one of the following tags [I-LOC B-ORG O B-PER I-PER I-MISC B-MISC I-ORG B-LOC]

```
import os
import torch
import torch.nn as nn
from torch.utils.data import Dataset, DataLoader

from transformers import (
    BertTokenizerFast,
    BertModel,
    get_scheduler
)


from sequeval.metrics import classification_report, f1_score, precision_score, recall_score

class CFG:
    model_name = "bert-base-cased"
    max_len = 128
    train_batch_size = 16
    valid_batch_size = 16
    lr = 5e-5
    epochs = 2
    device = "cuda" if torch.cuda.is_available() else "cpu"


# Label mapping (CoNLL2003)
label_list = ["O", "B-PER", "I-PER", "B-ORG", "I-ORG",
              "B-LOC", "I-LOC", "B-MISC", "I-MISC"]

label2id = {label: idx for idx, label in enumerate(label_list)}
id2label = {idx: label for label, idx in label2id.items()}

tokenizer = BertTokenizerFast.from_pretrained(CFG.model_name)
```

 /usr/local/lib/python3.12/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (<https://huggingface.co/settings/tokens>)
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public model

```
warnings.warn(
tokenizer_config.json: 100% 49.0/49.0 [00:00<00:00, 5.10kB/s]
vocab.txt: 100% 213k/213k [00:00<00:00, 7.34MB/s]
tokenizer.json: 100% 436k/436k [00:00<00:00, 34.4MB/s]
config.json: 100% 570/570 [00:00<00:00, 55.2kB/s]
```

```
# =====
#  Cell 4: Dataset loader
# =====
```

```
def read_conll_data(filepath):
    texts, tags = [], []
    words, labels = [], []
    with open(filepath, encoding="utf-8") as f:
        for line in f:
            if line.startswith("-DOCSTART-") or line.strip() == "":
                if words:
                    texts.append(words)
                    tags.append(labels)
                    words, labels = [], []
                continue
            splits = line.strip().split()
            words.append(splits[0])
            labels.append(splits[-1])
        if words: # last sentence
            texts.append(words)
            tags.append(labels)
    return texts, tags

train_texts, train_tags = read_conll_data("/content/train.txt")
valid_texts, valid_tags = read_conll_data("/content/valid.txt")
test_texts, test_tags = read_conll_data("/content/valid.txt")

print("✅ Loaded dataset:", len(train_texts), len(valid_texts), len(test_texts))
```

 ✅ Loaded dataset: 14041 3250 3250

```
class NERDataset(Dataset):
    def __init__(self, texts, tags, tokenizer, label2id, max_len):
        self.encodings = encode_examples(texts, tags, tokenizer, label2id, max_len)

    def __len__(self):
```

```

return len(self.encodings["input_ids"])

def __getitem__(self, idx):
    item = {key: torch.tensor(val[idx]) for key, val in self.encodings.items()}
    return item

train_dataset = NERDataset(train_texts, train_tags, tokenizer, label2id, CFG.max_len)
valid_dataset = NERDataset(valid_texts, valid_tags, tokenizer, label2id, CFG.max_len)

train_loader = DataLoader(train_dataset, batch_size=CFG.train_batch_size, shuffle=True)
valid_loader = DataLoader(valid_dataset, batch_size=CFG.valid_batch_size)

class NERModel(nn.Module):
    def __init__(self, model_name, num_labels):
        super(NERModel, self).__init__()
        self.bert = BertModel.from_pretrained(model_name)
        self.dropout = nn.Dropout(0.3)
        self.classifier = nn.Linear(self.bert.config.hidden_size, num_labels)

    def forward(self, input_ids, attention_mask, labels=None):
        outputs = self.bert(input_ids=input_ids, attention_mask=attention_mask)
        sequence_output = self.dropout(outputs.last_hidden_state)
        logits = self.classifier(sequence_output)

        loss = None
        if labels is not None:
            loss_fct = nn.CrossEntropyLoss(ignore_index=-100)
            loss = loss_fct(logits.view(-1, logits.shape[-1]), labels.view(-1))
        return loss, logits

model = NERModel(CFG.model_name, len(label2id)).to(CFG.device)

```



model.safetensors: 100%

436M/436M [00:12<00:00, 44.4MB/s]

```

optimizer = torch.optim.AdamW(model.parameters(), lr=CFG.lr)
num_training_steps = len(train_loader) * CFG.epochs
scheduler = get_scheduler("linear", optimizer, num_warmup_steps=0, num_training_steps=num_tr

```

```

def train_fn(model, dataloader):
    model.train()
    total_loss = 0
    for batch in dataloader:
        optimizer.zero_grad()
        input_ids = batch["input_ids"].to(CFG.device)
        attention_mask = batch["attention_mask"].to(CFG.device)
        labels = batch["labels"].to(CFG.device)

        loss, _ = model(input_ids, attention_mask, labels)

```

```

        loss.backward()
        optimizer.step()
        scheduler.step()
        total_loss += loss.item()
    return total_loss / len(dataloader)

def eval_fn(model, dataloader):
    model.eval()
    preds, true_labels = [], []
    total_loss = 0
    with torch.no_grad():
        for batch in dataloader:
            input_ids = batch["input_ids"].to(CFG.device)
            attention_mask = batch["attention_mask"].to(CFG.device)
            labels = batch["labels"].to(CFG.device)

            loss, logits = model(input_ids, attention_mask, labels)
            total_loss += loss.item()
            predictions = torch.argmax(logits, dim=-1)

            for i in range(len(labels)):
                pred_labels, true_l = [], []
                for j in range(len(labels[i])):
                    if labels[i][j] != -100:
                        pred_labels.append(id2label[predictions[i][j].item()])
                        true_l.append(id2label[labels[i][j].item()])
                preds.append(pred_labels)
                true_labels.append(true_l)

    f1 = f1_score(true_labels, preds)
    return total_loss / len(dataloader), f1, classification_report(true_labels, preds)

train_losses = []
valid_losses = []
valid_f1_scores = []

for epoch in range(CFG.epochs):
    train_loss = train_fn(model, train_loader)
    valid_loss, f1, report = eval_fn(model, valid_loader)
    train_losses.append(train_loss)
    valid_losses.append(valid_loss)
    valid_f1_scores.append(f1)
    print(f"Epoch {epoch+1}/{CFG.epochs} | Train Loss: {train_loss:.4f} | Valid Loss: {valid_loss:.4f} | F1: {f1:.4f}")
    print(report)

```



Epoch 1/2 | Train Loss: 0.0919 | F1: 0.9319

	precision	recall	f1-score	support
LOC	0.96	0.95	0.96	1837
MISC	0.84	0.88	0.86	922
ORG	0.90	0.91	0.91	1341
PER	0.96	0.97	0.96	1836

micro avg	0.93	0.94	0.93	5936
macro avg	0.91	0.93	0.92	5936
weighted avg	0.93	0.94	0.93	5936

Epoch 2/2 | Train Loss: 0.0219 | F1: 0.9462

	precision	recall	f1-score	support
LOC	0.97	0.96	0.97	1837
MISC	0.88	0.91	0.89	922
ORG	0.91	0.93	0.92	1341
PER	0.97	0.97	0.97	1836

micro avg	0.94	0.95	0.95	5936
macro avg	0.93	0.94	0.94	5936
weighted avg	0.94	0.95	0.95	5936

```
def predict_sentence(sentence, tokenizer, model, id2label, max_len=128):
    model.eval()
    tokens = sentence.split()

    encoding = tokenizer(
        tokens,
        is_split_into_words=True,
        return_tensors="pt",
        truncation=True,
        padding="max_length",
        max_length=max_len
    )
    word_ids = encoding.word_ids() # lấy trước khi move to device
    device = next(model.parameters()).device
    input_ids = encoding["input_ids"].to(device)
    attention_mask = encoding["attention_mask"].to(device)

    with torch.no_grad():
        _, logits = model(input_ids=input_ids, attention_mask=attention_mask)

    preds = logits.argmax(-1).squeeze(0).tolist()

    results = []
    seen_word = set()
    for wid, pred in zip(word_ids, preds):
        if wid is None:
            continue
        if wid in seen_word:
            # chỉ lấy subword đầu tiên
            continue
        seen_word.add(wid)
        results.append((tokens[wid], id2label[pred]))
    return results
```

```

sentences = [
    "Elon Musk founded SpaceX in the United States",
    "Cristiano Ronaldo joined Al Nassr in Saudi Arabia",
    "Apple Inc. is headquartered in Cupertino",
    "The Olympics will be held in Paris in 2024",
    "Barack Obama met Angela Merkel in Berlin",
    "Vietnam won the gold medal in the SEA Games",
    "Google and Microsoft are competing in artificial intelligence"
]

```

```

for s in sentences:
    print(s)
    print(predict_sentence(s, tokenizer, model, id2label))
    print("-----")

```

```

➡ Elon Musk founded SpaceX in the United States
[('Elon', 'B-PER'), ('Musk', 'I-PER'), ('founded', 'O'), ('SpaceX', 'B-ORG'), ('in', 'O')
-----
Cristiano Ronaldo joined Al Nassr in Saudi Arabia
[('Cristiano', 'B-PER'), ('Ronaldo', 'I-PER'), ('joined', 'O'), ('Al', 'B-ORG'), ('Nassr
-----
Apple Inc. is headquartered in Cupertino
[('Apple', 'B-ORG'), ('Inc.', 'I-ORG'), ('is', 'O'), ('headquartered', 'O'), ('in', 'O')
-----
The Olympics will be held in Paris in 2024
[('The', 'O'), ('Olympics', 'B-MISC'), ('will', 'O'), ('be', 'O'), ('held', 'O'), ('in',
-----
Barack Obama met Angela Merkel in Berlin
[('Barack', 'B-PER'), ('Obama', 'I-PER'), ('met', 'O'), ('Angela', 'B-PER'), ('Merkel',
-----
Vietnam won the gold medal in the SEA Games
[('Vietnam', 'B-LOC'), ('won', 'O'), ('the', 'O'), ('gold', 'O'), ('medal', 'O'), ('in',
-----
Google and Microsoft are competing in artificial intelligence
[('Google', 'B-ORG'), ('and', 'O'), ('Microsoft', 'B-ORG'), ('are', 'O'), ('competing',
-----

```

```

# Save the model
torch.save(model.state_dict(), "ner_model.pth")
print("Model saved successfully!")

```

```

➡ Model saved successfully!

```

✓ Named_Entity_Recognition_rule_based

```

import spacy
from spacy.tokens import Span
from spacy.language import Language
from spacy import displacy

```



```

import pandas as pd
import re
from collections import defaultdict
from tqdm import tqdm

@Language.component("custom_rule_based_ner")
def custom_rule_based_ner(doc):
    entities = []
    for token in doc:
        if token.is_title and token.text not in spacy.lang.en.stop_words.STOP_WORDS and not
            if token.i + 1 < len(doc) and doc[token.i + 1].is_title:
                span = Span(doc, token.i, token.i + 2, label="ORG")
            else:
                span = Span(doc, token.i, token.i + 1, label="PER")
            entities.append(span)

    org_pattern = r'\b[A-Z][a-z]+ (Inc\.|Corp\.|[A-Z]{2,})\b'
    for match in re.finditer(org_pattern, doc.text):
        start_char, end_char = match.span()
        start_token = min([t.i for t in doc if t.idx >= start_char], default=None)
        end_token = min([t.i for t in doc if t.idx + len(t.text) > end_char], default=None)
        if start_token is not None and end_token is not None and start_token < end_token:
            span = Span(doc, start_token, end_token, label="ORG")
            entities.append(span)

    for token in doc:
        if token.is_title and token.text not in [ent.text for ent in entities] and not token
            span = Span(doc, token.i, token.i + 1, label="MISC")
            entities.append(span)

    for token in doc:
        if token.like_num and not token.is_stop or token.text in spacy.lang.en.stop_words.ST
            span = Span(doc, token.i, token.i + 1, label="O")
            entities.append(span)

    # filter overlapping entities
    filtered_entities = []
    sorted_entities = sorted(entities, key=lambda ent: (ent.start , -ent.end))
    last_end = -1
    for ent in sorted_entities:
        if ent.start >= last_end:
            filtered_entities.append(ent)
            last_end = ent.end
    doc.ents = filtered_entities
    return doc

def process_text(nlp,token,modename):
    text = ' '.join(token)
    doc = nlp(text)

```

```

entities = [(ent.text, ent.label_) for ent in doc.ents]
df = pd.DataFrame(entities, columns=['Entity', 'Label'])
print(f"\nResults for {modename}:\n")
print(df)
return doc ,df

def visualize_entities(doc, model_name, sentence_text):
    print(f"\nVisualization for {model_name} on sentence: {sentence_text}")
    displacy.render(doc, style="ent", jupyter=False)
    print("\n" + "="*50 + "\n")

def evaluate_ner(predicted_df, true_labels, tokens):
    predicted_entities = set((row['Entity'], row['Label']) for _, row in predicted_df.iterrows())
    true_entities = set()
    i = 0
    while i < len(tokens):
        if true_labels[i].startswith('B-'):
            label = true_labels[i][2:]
            entity_tokens = [tokens[i]]
            j = i + 1
            while j < len(tokens) and true_labels[j].startswith('I-') and true_labels[j][2:] == label:
                entity_tokens.append(tokens[j])
                j += 1
            entity_text = ' '.join(entity_tokens)
            true_entities.add((entity_text, label))
            i = j
        else:
            i += 1

    true_positive = len(predicted_entities & true_entities)
    false_positive = len(predicted_entities - true_entities)
    false_negative = len(true_entities - predicted_entities)

    precision = true_positive / (true_positive + false_positive) if (true_positive + false_positive) > 0 else 0
    recall = true_positive / (true_positive + false_negative) if (true_positive + false_negative) > 0 else 0
    f1_score = 2 * (precision * recall) / (precision + recall) if (precision + recall) > 0 else 0

    print(f"Precision: {precision:.2f}, Recall: {recall:.2f}, F1-Score: {f1_score:.2f}")

    return precision, recall, f1_score

def process_dataset(nlp_rule, nlp_sm, nlp_lg, dataset, dataset_name, sample_sentence=None):
    metrics = defaultdict(list)
    sample_processed = False

    if not dataset:
        print(f"Warning: {dataset_name} dataset is empty")
        return metrics

```

```

for i, (tokens, true_labels) in tqdm(enumerate(dataset[:10]), total=min(10, len(dataset)))
    if not tokens or not true_labels:
        print(f"Skipping empty or invalid sentence {i} in {dataset_name}")
        continue

    text = " ".join(tokens)
    if sample_sentence and text.startswith(sample_sentence) and not sample_processed:
        print(f"\nProcessing sample sentence from {dataset_name}: {text}")
        print(f"Ground truth labels: {true_labels}")
        sample_processed = True

# Rule-based NER
try:
    doc_rule, df_rule = process_text(nlp_rule, tokens, "Rule-Based Model")
    precision_rule, recall_rule, f1_rule = evaluate_ner(df_rule, true_labels, tokens)
    metrics['rule_precision'].append(precision_rule)
    metrics['rule_recall'].append(recall_rule)
    metrics['rule_f1'].append(f1_rule)
except Exception as e:
    print(f"Error processing Rule-Based Model for sentence {i} in {dataset_name}: {e}")
    continue

# Small model NER
try:
    doc_sm, df_sm = process_text(nlp_sm, tokens, "Small Model")
    precision_sm, recall_sm, f1_sm = evaluate_ner(df_sm, true_labels, tokens)
    metrics['sm_precision'].append(precision_sm)
    metrics['sm_recall'].append(recall_sm)
    metrics['sm_f1'].append(f1_sm)
except Exception as e:
    print(f"Error processing Small Model for sentence {i} in {dataset_name}: {e}")
    continue

# Large model NER
try:
    doc_lg, df_lg = process_text(nlp_lg, tokens, "Large Model")
    precision_lg, recall_lg, f1_lg = evaluate_ner(df_lg, true_labels, tokens)
    metrics['lg_precision'].append(precision_lg)
    metrics['lg_recall'].append(recall_lg)
    metrics['lg_f1'].append(f1_lg)
except Exception as e:
    print(f"Error processing Large Model for sentence {i} in {dataset_name}: {e}")
    continue

if sample_sentence and text.startswith(sample_sentence):
    visualize_entities(doc_rule, "Rule-Based Model", text)
    visualize_entities(doc_sm, "Small Model", text)
    visualize_entities(doc_lg, "Large Model", text)
    print("\nComparison for Sample Sentence:")
    print("\nRule-Based Entities:")
    print(df_rule[['Entity', 'Label']])

```

```

print("\nSmall Model Entities:")
print(df_sm[['Entity', 'Label']])
print("\nLarge Model Entities:")
print(df_lg[['Entity', 'Label']])

if not metrics['rule_precision']:
    print(f"Warning: No metrics collected for {dataset_name}")

summary = pd.DataFrame({
    'Model': ['Rule-Based', 'Small Model', 'Large Model'],
    'Precision': [sum(metrics['rule_precision']) / len(metrics['rule_precision']) if metrics['rule_precision'] else 0,
                  sum(metrics['sm_precision']) / len(metrics['sm_precision']) if metrics['sm_precision'] else 0,
                  sum(metrics['lg_precision']) / len(metrics['lg_precision']) if metrics['lg_precision'] else 0],
    'Recall': [sum(metrics['rule_recall']) / len(metrics['rule_recall']) if metrics['rule_recall'] else 0,
               sum(metrics['sm_recall']) / len(metrics['sm_recall']) if metrics['sm_recall'] else 0,
               sum(metrics['lg_recall']) / len(metrics['lg_recall']) if metrics['lg_recall'] else 0],
    'F1-Score': [sum(metrics['rule_f1']) / len(metrics['rule_f1']) if metrics['rule_f1'] else 0,
                  sum(metrics['sm_f1']) / len(metrics['sm_f1']) if metrics['sm_f1'] else 0,
                  sum(metrics['lg_f1']) / len(metrics['lg_f1']) if metrics['lg_f1'] else 0],
})
print(f"\nSummary for {dataset_name}:\n")
print(summary)

return metrics

```

! python -m spacy download en_core_web_lg



Collecting en-core-web-lg==3.8.0

Downloading https://github.com/explosion/spacy-models/releases/download/en_core_web_lg-3.8.0/en_core_web_lg-3.8.0.tar.gz

400.7/400.7 MB 3.9 MB/s eta 0:00:00

Installing collected packages: en-core-web-lg

Successfully installed en-core-web-lg-3.8.0

✓ Download and installation successful

You can now load the package via `spacy.load('en_core_web_lg')`

⚠ Restart to reload dependencies

If you are in a Jupyter or Colab notebook, you may need to restart Python in order to load all the package's dependencies. You can do this by selecting the 'Restart kernel' or 'Restart runtime' option.

```

file_paths = {
    "train": "/content/train.txt",
    "test": "/content/test.txt",
    "valid": "/content/valid.txt"
}

data_sets = {name : read_conll_data(path) for name, path in file_paths.items()}

nlp_small = spacy.load("en_core_web_sm", disable=['lemmatizer', 'textcat'])
nlp_large = spacy.load("en_core_web_lg", disable=['lemmatizer', 'textcat'])

```

```

nlp_rule = spacy.blank("en")
ner = nlp_rule.add_pipe("ner")

for label in ["PER", "ORG", "MISC", "O", "LOC"]:
    ner.add_label(label)

nlp_rule.add_pipe("custom_rule_based_ner")
nlp_rule.initialize()

nlp_rule.max_length = 2000000
nlp_small.max_length = 2000000
nlp_large.max_length = 2000000

sample_sentence = "The European Commission said on Thursday"

all_metrics = {}
for name, dataset in data_sets.items():
    print(f"\nProcessing {name} dataset...")
    texts, tags = dataset # Unpack the dataset tuple into texts and tags
    # Iterate over zipped texts and tags
    metrics = process_dataset(nlp_rule, nlp_small, nlp_large, list(zip(texts, tags)), name,
    all_metrics[name] = metrics

print("\nCross-Dataset Comparison:")
metric_mapping = {
    'Precision': 'precision',
    'Recall': 'recall',
    'F1-Score': 'f1'
}
for metric in ['Precision', 'Recall', 'F1-Score']:
    print(f"\nAverage {metric}:")
    comparison = pd.DataFrame({
        "Model": ["Rule-Based", "en_core_web_sm", "en_core_web_lg"],
        "Train": [
            sum(all_metrics['train'][f'rule_{metric_mapping[metric]}']) / len(all_metrics['train'])
            sum(all_metrics['train'][f'sm_{metric_mapping[metric]}']) / len(all_metrics['train'])
            sum(all_metrics['train'][f'lg_{metric_mapping[metric]}']) / len(all_metrics['train'])
        ],
        "Valid": [
            sum(all_metrics['valid'][f'rule_{metric_mapping[metric]}']) / len(all_metrics['valid'])
            sum(all_metrics['valid'][f'sm_{metric_mapping[metric]}']) / len(all_metrics['valid'])
            sum(all_metrics['valid'][f'lg_{metric_mapping[metric]}']) / len(all_metrics['valid'])
        ],
        "Test": [
            sum(all_metrics['test'][f'rule_{metric_mapping[metric]}']) / len(all_metrics['test'])
            sum(all_metrics['test'][f'sm_{metric_mapping[metric]}']) / len(all_metrics['test'])
            sum(all_metrics['test'][f'lg_{metric_mapping[metric]}']) / len(all_metrics['test'])
        ]
    })
    print(comparison)

```



Precision: 0.00, Recall: 0.00, F1-Score: 0.00

```
# Save the spaCy models
nlp_small.to_disk("en_core_web_sm_saved")
nlp_large.to_disk("en_core_web_lg_saved")
print("spaCy models saved successfully!")
```

🔄 spaCy models saved successfully!

Start coding or [generate](#) with AI.

✓ Task

Create a Gradio interface to compare the NER predictions of a BERT model and two spaCy models (en_core_web_sm and en_core_web_lg).

✓ Install gradio

Subtask:

Add a code cell to install the Gradio library.

Reasoning: The subtask is to install the Gradio library. This requires adding a new code cell with the pip install command for gradio.

```
!pip install gradio
```

🔄 Requirement already satisfied: gradio in /usr/local/lib/python3.12/dist-packages (5.43.1)
 Requirement already satisfied: aiofiles<25.0,>=22.0 in /usr/local/lib/python3.12/dist-packages (22.0)
 Requirement already satisfied: anyio<5.0,>=3.0 in /usr/local/lib/python3.12/dist-packages (3.0)
 Requirement already satisfied: brotli>=1.1.0 in /usr/local/lib/python3.12/dist-packages (1.1.0)
 Requirement already satisfied: fastapi<1.0,>=0.115.2 in /usr/local/lib/python3.12/dist-packages (0.115.2)
 Requirement already satisfied: ffmpy in /usr/local/lib/python3.12/dist-packages (from gradio) (0.3.2)
 Requirement already satisfied: gradio-client==1.12.1 in /usr/local/lib/python3.12/dist-packages (1.12.1)
 Requirement already satisfied: groovy~=0.1 in /usr/local/lib/python3.12/dist-packages (0.1)
 Requirement already satisfied: httpx<1.0,>=0.24.1 in /usr/local/lib/python3.12/dist-packages (0.24.1)
 Requirement already satisfied: huggingface-hub<1.0,>=0.33.5 in /usr/local/lib/python3.12/dist-packages (0.33.5)
 Requirement already satisfied: jinja2<4.0 in /usr/local/lib/python3.12/dist-packages (from gradio) (3.1.2)
 Requirement already satisfied: markupsafe<4.0,>=2.0 in /usr/local/lib/python3.12/dist-packages (2.1.5)
 Requirement already satisfied: numpy<3.0,>=1.0 in /usr/local/lib/python3.12/dist-packages (1.26.4)
 Requirement already satisfied: orjson~=3.0 in /usr/local/lib/python3.12/dist-packages (3.10.11)
 Requirement already satisfied: packaging in /usr/local/lib/python3.12/dist-packages (from gradio) (24.1)
 Requirement already satisfied: pandas<3.0,>=1.0 in /usr/local/lib/python3.12/dist-packages (2.2.3)
 Requirement already satisfied: pillow<12.0,>=8.0 in /usr/local/lib/python3.12/dist-packages (10.4.0)

```

Requirement already satisfied: pydantic<2.12,>=2.0 in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: pydub in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: python-multipart>=0.0.18 in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: pyyaml<7.0,>=5.0 in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: ruff>=0.9.3 in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: safehttpx<0.2.0,>=0.1.6 in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: semantic-version~=2.0 in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: starlette<1.0,>=0.40.0 in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: tomlkit<0.14.0,>=0.12.0 in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: typer<1.0,>=0.12 in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: typing-extensions~=4.0 in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: uvicorn>=0.14.0 in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: fsspec in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: websockets<16.0,>=10.0 in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: idna>=2.8 in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: sniffio>=1.1 in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: certifi in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: httpcore==1.* in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: h11>=0.16 in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: filelock in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: requests in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: tqdm>=4.42.1 in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: hf-xet<2.0.0,>=1.1.3 in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: annotated-types>=0.6.0 in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: pydantic-core==2.33.2 in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: typing-inspection>=0.4.0 in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: click>=8.0.0 in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: shellingham>=1.3.0 in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: rich>=10.11.0 in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: markdown-it-py>=2.2.0 in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: charset_normalizer<4,>=2 in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)
Requirement already satisfied: mdurl~=0.1 in /usr/local/lib/python3.12/dist-packages (from gradio==4.43.0)

```

✓ Load models

Subtask:

Load the saved BERT model and the saved spaCy models (en_core_web_sm_saved and en_core_web_lg_saved).

Reasoning: Load the saved BERT and spaCy models for use in the Gradio interface.

```

model.load_state_dict(torch.load("ner_model.pth", map_location=CFG.device))
nlp_small_loaded = spacy.load("en_core_web_sm_saved")

```



```
nlp_large_loaded = spacy.load("en_core_web_lg_saved")
```

```
print("Models loaded successfully!")
```

```
⇒ Models loaded successfully!
```

✓ Define prediction functions

Subtask:

Create Python functions that take a sentence as input and return NER predictions in a user-friendly format for the BERT model and both spaCy models.

Reasoning: Define the functions to get NER predictions from the BERT and spaCy models.

```
def get_bert_predictions(sentence, tokenizer, model, id2label, max_len=128):
    """
    Gets NER predictions from the BERT model for a given sentence.

    Args:
        sentence (str): The input sentence.
        tokenizer: The BERT tokenizer.
        model: The trained BERT model.
        id2label (dict): Dictionary mapping label IDs to labels.
        max_len (int): Maximum sequence length.

    Returns:
        list: A list of (word, label) tuples.
    """
    results = predict_sentence(sentence, tokenizer, model, id2label, max_len)
    return results


def get_spacy_predictions(nlp_model, sentence):
    """
    Gets NER predictions from a spaCy model for a given sentence.

    Args:
        nlp_model: The loaded spaCy language model.
        sentence (str): The input sentence.

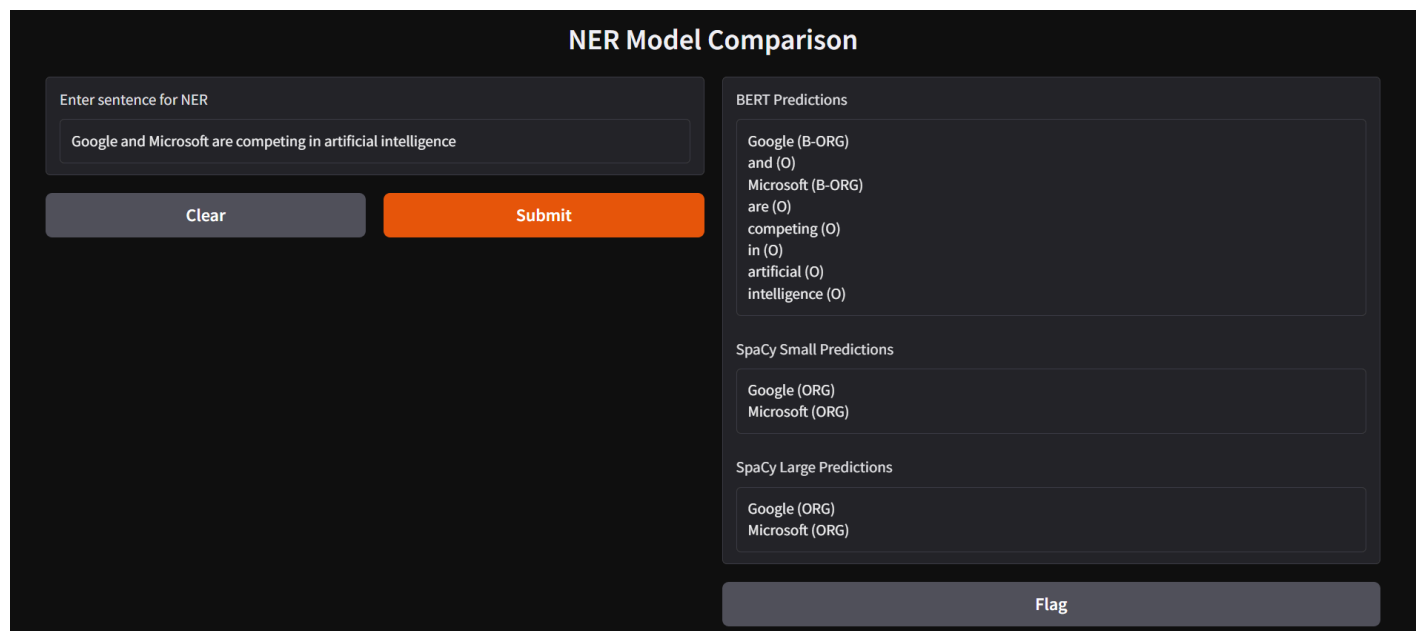
    Returns:
        list: A list of (entity_text, label) tuples.
    """
    doc = nlp_model(sentence)
    entities = [(ent.text, ent.label_) for ent in doc.ents]
    return entities
```

✓ Build gradio interface

Subtask:

Design and implement a Gradio interface with input text boxes for the sentence and output displays for the predictions from each model.

Reasoning: Implement the `ner_comparison` function to call the prediction functions and format the output, then create the Gradio interface.



The image shows a Gradio interface titled "NER Model Comparison". It features a text input box on the left with the placeholder "Enter sentence for NER" and the text "Google and Microsoft are competing in artificial intelligence". Below the input box are two buttons: "Clear" and "Submit". To the right of the input box, there are three output boxes for predictions: "BERT Predictions", "SpaCy Small Predictions", and "SpaCy Large Predictions". The BERT Predictions box shows the following output: "Google (B-ORG)", "and (O)", "Microsoft (B-ORG)", "are (O)", "competing (O)", "in (O)", "artificial (O)", "intelligence (O)". The SpaCy Small Predictions box shows: "Google (ORG)", "Microsoft (ORG)". The SpaCy Large Predictions box shows: "Google (ORG)", "Microsoft (ORG)". At the bottom right of the interface is a "Flag" button.

Model	Predictions
BERT	Google (B-ORG), and (O), Microsoft (B-ORG), are (O), competing (O), in (O), artificial (O), intelligence (O)
SpaCy Small	Google (ORG), Microsoft (ORG)
SpaCy Large	Google (ORG), Microsoft (ORG)

... The End!